

Computação Paralela

October-2020

Program parallelization with OpenMP (part2)

These exercises aim to introduce concepts of program parallelization for shared memory platforms with OpenMP. Compile the program with `gcc -O3 -std=c99 -fopenmp -lm <<file>>`. Set the number of threads before running the program with `export OMP_NUM_THREADS=<<NT>>`

- 1) Consider the following OpenMP program to compute the dot product of two vectors

```
#define size 100000
double a[size], b[size];

void main() {
    // init
    for(int i=0; i<size; i++) {
        a[i] = 1.0/((double) (size-i));
        b[i] = a[i] * a[i];
    }

    // computedot
    double dot = 0;

    for(int i=0; i<size; i++) {
        dot += a[i]*b[i];
    }

    printf("Dot is %18.16f\n", dot);
}
```

- a) Parallelize the code by including the *omp parallel for* pragma in the second loop (*computedot*). Are you able to get same results with different number of threads and in different runs? Why? (suggestion: is the computed value higher or lower than expected? How does it changes with the number of threads?). What is needed for correct computation?
- b) Use the *reduction* clause to compute the dot correctly. Compile and run your code also without OpenMP. Do you get exactly same results in all cases?
- 2) Develop a parallel version of the following program that computes an approximation of the value of pi (identify where is the data race). Measure the program speed-up by increasing the number of threads. Use the *time ./a.out* command to measure the execution time.

```
double f( double a ) {
    return (4.0 / (1.0 + a*a));
}

double pi = 3.141592653589793238462643;

int main() {
    double mypi = 0;
    int n = 100000000; // number of points to compute
    float h = 1.0 / n;

    for(int i=0; i<n; i++) {
        mypi = mypi + f(i*h);
    }
    mypi = mypi * h;
    printf(" pi = %.10f \n", (pi - mypi));
}
```